

Embracing the new threat: towards automatically, self-diversifying malware

Mathias Payer

mathias.payer@nebelwelt.net

UC Berkeley

Abstract

Signature-based similarity metrics are the primary mechanism to detect malware on current systems. Each file is scanned and compared against a set of signatures. This approach has several problems: (i) all possible detectable malware must have a signature in the database and (ii) it might take a substantial amount of time between initial spread of the malware and the time anti-malware companies generate a signature to protect from the malware.

On the other hand, the malware landscape is changing: there are only few malware families alive at a certain point in time. Each family evolves along a common software update and maintenance cycle. Individual malware instances are repacked or obfuscated whenever they are detected by a large set of anti-malware products, basically resulting in an arms race between malware authors and anti-malware products.

Anti-malware products are not efficient if they follow this arms race and we show how it is possible to maximize the advantage for malware distributors. We present MalDiv, an automatic diversification mechanism that uses compiler-based transformations to generate an almost infinite amount of binaries with the same functionality but very low similarity, resulting in different signatures. Malware diversity builds on software diversity and uses open decisions in the compiler to reorder and change code and data. In addition, static data is encrypted using a set of transformations. Such a tool allows malware distributors to generate an almost unlimited amount of binaries that cannot be detected using signature-based matching.

1. Introduction

The malware landscape is changing but current defenses still rely on assumptions made more than 20 years ago: reaction and adaptation to new malware families is slow and the amount of different new malware samples an anti-malware company can cope with is limited due to some manual analysis involved in the process of generating new signatures. On the malware side, it is all about money. Attacks are (usually) no longer carried out by hobbyists or home grown hackers, but by larger groups that adhere to a strict software develop-

ment cycle. A new piece of malware is developed over some time, sold to different customers, and updated in a regular software release maintenance and development cycle.

There are no longer thousands of different viruses, worms, and other malware alive at any point in time but only a few malware families. These few malware families are updated on a regular basis and diversified using packers (or crypters) [15, 23, 24, 26, 28]. Packers are small programs that hide the real code and data of the malware and unpack (or decrypt) it at runtime. The advantage of using a packer is that only a small amount of code must be changed to produce a completely different binary that evades signature-based detection as carried out by current defense technologies like anti-malware products. Yet, each new version incurs some cost. Somebody needs to manually update the packer or update the configuration of the packer to produce a new, manually diversified version of a specific malware family.

Due to the huge amount of manually diversified and packed binaries of only few malware families found in the wild we infer that malware diversification is already used in practice to circumvent signature-based matching, albeit using a slow, partly manual process.

Classic defenses are severely limited when facing diversified binaries. Anti-malware products (basically all anti-virus products evolved into anti-malware products) still work according to old assumptions of anti-virus software made in the late 80ies or early 90ies: each file is scanned and compared against a database of well-known bad examples. This database must be updated regularly and deployed to all the clients. A slow update cycle allows new malware to spread fast and far before it can be detected. In addition, the resources of defensive systems are constraint in multiple ways: (i) the amount of performance slowdown a user accepts is limited, (ii) the update frequency of the database is limited by the amount of bandwidth and server resources that the anti-malware company has available, and (iii) the number of malware samples that can be analyzed is constrained by the amount of reverse engineers that the company can employ.

Defending current systems against malware is already a hard task (relying on these old practices). In this paper we present MalDiv, a new way to automatically diversify

malware using a compiler-based diversification approach. This approach allows us to quickly and automatically generate large amounts of unique binaries using a single, unique source code base. MalDiv produces binaries with low similarity between each other using the same source code. MalDiv is effective, if two binaries produced from the same source code have low similarity, e.g., a signature for a first file does not match a second file.

For MalDiv we rely on software diversity [4] as a baseline. Software diversity uses open choices during the compilation of source code to accomplish different goals, e.g., it can be used to randomize the data layout on the heap to protect against memory corruption (an exploit only succeeds on the specific diversified binary it was written for but not on other diversified binaries) [4, 5, 9, 10, 13, 17]. The goal of malware diversity is to minimize similarity between different binaries and we diversify both code and data.

If the malware authors can produce new (different) binaries faster than the anti-malware companies can come up with signatures then signature-based detection approaches are completely mitigated. In addition, MalDiv can be tuned to mitigate against other similarity metrics and heuristics as well, resulting in an arms race between the anti-malware companies that need to come up with new (more sophisticated and probably more resource intensive) matching techniques.

The paper is organized as follows: Section 2 introduces background information in malware detection (presenting different techniques used to detect malware); Section 3 introduces MalDiv, our software diversity mechanism geared towards malware diversity; Section 4 evaluates MalDiv; and Section 5 concludes.

2. Malware detection

On a high level each anti-malware product offers comparable functionality: it (i) scans each file before it is executed or used; (ii) extracts a signature of the scanned file; (iii) compares the signature against a database of known bad signatures (using some similarity metric); and (iv) either puts the file into quarantine if it matches or allows execution or usage of the file. Anti-malware products differ in the size of the database, the speed of their response to newly diversified malware samples, and their heuristics for some similarity metrics.

The signature database is updated frequently and anti-malware companies collect large sets of potentially malicious binary samples. These samples are analyzed first automatically and later manually to generate precise signatures if they are malicious. This manual analysis limits the amount of signatures that can be generated in a specific amount of time.

Anti-malware products run under the assumption that they are installed on a clean and uncompromised machine and that they will catch all malware. If any malware is al-

lowed to execute (on the same or higher privilege level as the anti-malware software) then all security guarantees given by the anti-malware product are void (as the malware might disable or remove the anti-malware software). So in the end, anti-malware software uses similar techniques like malware to integrate itself deeply into the operating system (or more recently the hypervisor) to protect itself from malware access. Unfortunately, users are only willing to accept limited performance impact so anti-malware software must scan all files using a limited computational budget.

Modern anti-malware products, e.g., ClamAV [18], a well-known open-source anti-malware scanner rely on a set of similarity metrics with corresponding signatures:

Hash-based matching is a static technique that calculates a hash (e.g., MD5, or SHA1) of a file or part of a file and compares this simple hash against the signature for this similarity metric. This technique uses very low resources but the similarity metric fails if a single byte of the hashed area changes.

Sequence-based matching is a static technique that matches a small static sequence of bytes that identifies the malware family. This technique is more resilient against changes in the binary but the signature writer must be more careful not to match benign files as well.

Regular-expression-based matching is another static technique that uses a regular expression as signature, extending the computational expressiveness of sequence-based matching. A malware instance is identified if the regular expression matches a binary.

Behavioral matching is a static technique that identifies malware according to a set of heuristics, e.g., system calls commonly used in malware or other potentially malicious functionality. This technique is a general detection technique as it matches “general malicious behavior” and not specific malware instances.

Dynamic behavioral matching is a dynamic technique that extends behavioral matching. It uses a sandbox to analyze the potentially malicious files. All interactions with the operating system and modifications of the runtime system are monitored and matched against a set of heuristics. Advantages of this technique are that the effects of packers or other diversification can be mitigated but due to limited resources the analysis can only run for a short amount of time and is therefore incomplete.

All signature-based techniques have the disadvantage that they only match prior known malware that is in the signature database. Several other approaches have been developed to mitigate the drawbacks of signature-based matching. Some approaches detect the packer and try to unpack the actual malware code [15, 23, 26, 28]. Unfortunately, this approach only works for known packers or packing techniques. Other approaches try to normalize binaries to undo the effects of

diversification [3]. These techniques are limited to simple transformations (comparable to peephole optimizations carried out by the code generator in a compiler).

All behavioral techniques must cope with false positives and the incompleteness of their heuristics, resulting in an arms race with the attackers. Dynamic techniques face the additional drawback that the sandbox can be detected by the malware [2, 7, 11, 25, 27, 29] resulting in another arms race that tries to detect if a piece of code evades analysis [1, 7, 14, 16, 19, 21, 22]. Many of these dynamic techniques are too resource intensive to be used on consumer machines.

Other similarity metrics outside of anti-malware products can be used as well. Bindiff [8], for example, reconstructs the control-flow graph of the binary and uses graph-based matching as a similarity metric between two binaries.

3. Malware diversity

Malware diversity is a new form of software diversity that maximizes the differences between individual compiled instances to reduce the similarity. Compared to existing software diversification, malware diversity does not need to produce high performance code. Malware diversity disrupts similarity metrics by adding many small changes to code and data layout throughout the produced binary.

MalDiv diversifies both code and data. The code is diversified using already existing diversification techniques like instruction selection, register selection, control-flow changes, or side-effect free computation. Malware diversity focuses on maximum diversity in the generated code. Data is diversified using per-instance encodings of static data.

3.1 Software diversity

Software diversity [4] adds a different spin to compilation: under the assumption that the compiler version, flags, and source code remain static for all compilations software diversity produces different binaries for each compilation instead of deterministically producing the same binary over and over again. Software diversity can be used for many purposes: (i) to combat the “software monoculture” [10, 12, 13], (ii) to hide steganographic messages in binaries [6], or (iii) to protect software against exploits [4, 5, 9, 10, 17]. Diversification engines use different compiler techniques to produce different code and/or data layout for each compilation:

Instruction selection the code generator may select many different instructions to carry out some computation (e.g., an addition can be expressed using a subtraction) or may reorder instructions that do not depend on each other at will.

Register selection the register allocator may shuffle the list of free general purpose registers at any point in time shuffle or enforce an artificial sparsity of general purpose registers.

Control-flow changes the compiler can reorder, merge, or split basic blocks, reorder branch targets in switch statements, or use different inlining settings.

Side-effect free instructions the compiler can add garbage computation alongside the regular computation at will or even weave a second program into the program to change the instruction mix.

Variable diversification the compiler may reorder the variables on the stack. If the compiler can prove that no pointer ever escapes the diversified part of the program it can also reorder and modify the contents of structures and objects.

In general, software diversity relies on small changes of the aggressiveness in the optimization toolchain of a modern optimizing compiler. For example, if common subexpression elimination is less aggressive and only (randomly) removes some (random) subexpressions at some (random) points in time then the diversity is increased.

3.2 Data diversification

In addition to all the code diversification and data layout diversification, each binary also contains static data that is used during computation. This static data is often an ideal candidate for partial signatures. For MalDiv, we have looked into several diversification patterns for this static data. Of course, other diversification approaches for static data are possible and can easily be added to the diversification toolchain.

1. During compilation the static data is encrypted using a random key. At startup all data is decrypted before the original application starts.
2. A second approach compresses all static data into a small blob of data that unpacks the data at runtime.
3. A third approach decrypts the data only when it is first used, adding a short test if the data has already been decrypted before each use. This approach helps against dynamic behavioral analysis that tries to match the data region of the running malware process after unpacking the static data.

A large set of different encryption/decryption schemes are possible, from simple XOR-based encryption to more sophisticated encryption schemes. To protect the decryption code in the binary from becoming a signature it is diversified alongside the other source code.

3.3 Implementation

We implemented MalDiv, a prototype of our malware diversity engine on top of the multicompile [13], building on the LLVM [20] compiler framework. We use a combination of the existing diversity mechanisms and provide pre-set settings that maximize the diversity for the existing diversification mechanisms. In addition, we have added the simple XOR-based data diversification mechanism.

The project GitHub page <http://github.com/gannimo/MalDiv> contains all the resources needed to build the diversifying compiler, instructions on how to get the original multicompile, the static data diversification module, and a set of examples that show how to use the diversifying compiler.

4. Evaluation

This section presents an evaluation of the current MalDiv prototype implementation. We compare the similarity of a set of diversified binaries using different metrics: hash-based similarity and longest common substrings. We also discuss possible mitigations against malware diversity.

As binaries we use a set of SPEC CPU2006 benchmarks written in C/C++, nmap, and a simple port scanner.

4.1 Hash-based similarity

Hash-based similarity is not effective for any of the diversified binaries. Both the binary as a whole and individual segments (code and data regions) change due to the diversification between individual instances. Due to these changes neither whole-binary nor segment-based/partial hashing of binaries can be used to detect instances of the same malware.

4.2 Common substrings

To evaluate the power of individual signatures we test a set of diversified binaries for common substrings (a sequence of bytes that occurs in both files). Finding common substrings is computationally very expensive but allows us to check if some part of a binary remains static across several diversified instances.




Figure 1. List of common subsequences for SPEC CPU2006 benchmarks (log scale).

Figure 1 shows (i) the list of SPEC benchmarks always comparing two different diversified versions with each other, (ii) perlbench compared with bz2 and perlbench compared with gobmk to show the similarity between two different programs, and (iii) two diversified versions of nmap and a simple port scanner.

From Figure 1 we see that few shared substrings are longer than 20 bytes of length and most shared substrings are shorter than 20 bytes. The longer substrings all fall into one of the following criteria: start files (shared among all programs compiled by the same compiler), function call sequences (pushing parameters on to the stack), mov sequences (a set of mov instructions that initialize structures), floating point sequences that are currently not (yet) diversified, or hand written assembler instructions. Of these substrings only the handwritten assembly instructions pose a problem as they cannot be automatically broken up by the compiler. Apart from these limitations it will be hard for malware analysts to come up with efficient signatures for diversified binaries.

4.3 Possible mitigation

The evaluation showed that any signature-based matching is no longer effective for diversified binaries. Combined with the high frequency in which diversified binaries can be generated malware diversity allows malware authors to spread fast and far without wide-spread detection (individual malware instances can still be analyzed on a per-case basis).

Static and dynamic behavioral analysis of individual malware samples is of course possible and will discover simple malware families. We assume that malware authors will rely on existing anti-debugging and anti-forensics tools to mitigate these risks. Anti-debugging and anti-forensics is orthogonal to malware diversity and not part of MalDiv. Malware products will quickly incorporate anti-debugging and anti-forensic measurements on their own.

Graph-based binary similarity analysis tools like bindiff can still effectively recover some similarity between diversified instances due to the limited diversification of the control-flow that is implemented in the current MalDiv prototype. Bindiff shows reasonably high similarity between diversified instances and reasonably low similarity between functionally different binaries due to the differences in the control flow graph. Unfortunately, bindiff has a very high analysis cost of up to 15 minutes to compare two diversified instances. We reason that such a high analysis overhead will remain unpractical for wide-spread deployment of such an analysis technique.

5. Conclusion

Malware detection engines rely on signatures and common behavior to successfully detect malware. The assumption of these detection mechanisms is that each malware (family) can be classified using some form of signature.

In this paper we introduce malware diversity which breaks the above mentioned assumption. Our prototype MalDiv, a malware diversification technique automatically diversifies malware during the compilation. The LLVM-based technique produces a large amount of binaries with very low similarity in a short amount of time.

Acknowledgments

The idea to this project originates from a discussion between Stefan Brunthaler, Per Larsen, and Mathias Payer while Mathias visited UC Irvine. The work presented here (ab)uses all the great multicompile [13] work of Michael Franz’s group to maximize the differences between produced binaries to minimize similarity according to different metrics and adds a data diversification module for static data. This project builds on a lot of feedback, great discussions, and help from Michael Franz, Stefan Brunthaler, Per Larsen, Richard Wartell, and Stephen Crane.

References

- [1] D. Balzarotti, M. Cova, C. Karlberger, E. Kirda, C. Kruegel, and G. Vigna. Efficient detection of split personalities in malware. In *NDSS’10: Proc. Network and Distributed System Security Symp.*, 2010.
- [2] X. Chen, J. Andersen, Z. M. Mao, M. Bailey, and J. Nazario. Towards an Understanding of Anti-Virtualization and Anti-Debugging Behavior in Modern Malware. In *DSN’08: Proc. 38th Annual IEEE Int. Conf. on Dependable Systems and Networks*, pages 177–186, June 2008.
- [3] M. Christodorescu, J. Kinder, S. Jha, S. Katzenbeisser, and H. Veith. Malware normalization. Technical report, Technische Universität München, 2005.
- [4] F. B. Cohen. Operating system protection through program evolution. *Comput. Secur.*, 12(6):565–584, 1993. ISSN 0167-4048. URL [http://dx.doi.org/10.1016/0167-4048\(93\)90054-9](http://dx.doi.org/10.1016/0167-4048(93)90054-9).
- [5] B. Cox, D. Evans, A. Filipi, J. Rowanhill, W. Hu, J. Davidson, J. Knight, A. Nguyen-Tuong, and J. Hiser. N-variant systems: a secretless framework for security through diversity. In *Proceedings of the 15th conference on USENIX Security Symposium - Volume 15*, USENIX-SS’06, 2006. URL <http://dl.acm.org/citation.cfm?id=1267336.1267344>.
- [6] R. El-khalil and A. D. Keromytis. Hydan: Hiding information in program binaries. In *ICICS’04: Int. Conf. on Information and Communications Security*, 2004.
- [7] P. Ferrie. Attacks on virtual machine emulators. http://www.symantec.com/avcenter/reference/Virtual_Machine_Threats.pdf, 2006.
- [8] H. Flake. Structural comparison of executable objects. In *DIMVA’04*.
- [9] S. Forrest, A. Somayaji, and D. H. Ackley. Building diverse computer systems. In *Workshop on Hot Topics in Operating Systems*, pages 67–72, 1997.
- [10] M. Franz. E unibus pluram: massive-scale software diversity as a defense mechanism. In *Proc. 2010 workshop on New security paradigms*, NSPW ’10, pages 7–16, 2010. ISBN 978-1-4503-0415-3. URL <http://doi.acm.org/10.1145/1900546.1900550>.
- [11] T. Garfinkel, K. Adams, A. Warfield, and J. Franklin. Compatibility is not transparency: Vmm detection myths and realities. In *HOTOS’07: Proc. 11th USENIX workshop on Hot topics in operating systems*, pages 6:1–6:6, 2007. URL <http://dl.acm.org/citation.cfm?id=1361397.1361403>.
- [12] D. Geer, C. P. Pfleeger, B. Schneider, J. S. Quarterman, P. Metzger, R. Bace, and P. Gutmann. CyberInsecurity: The cost of monopoly – how the dominance of Microsoft’s products poses a risk to security. *Computer & Communications Industry Association Report*, 2003.
- [13] A. Homescu, S. Neisius, P. Larsen, S. Brunthaler, and M. Franz. Profile-guided automated software diversity. In *Proceedings of the 2013 IEEE/ACM International Symposium on Code Generation and Optimization*, CGO’13, 2013.
- [14] N. M. Johnson, J. Caballero, K. Z. Chen, S. McCamant, P. Poosankam, D. Reynaud, and D. Song. Differential slicing: Identifying causal execution differences for security applications. In *IEEE S&P’11*.
- [15] M. G. Kang, P. Poosankam, and H. Yin. Renovo: a hidden code extractor for packed executables. In *Proceedings of the 2007 ACM workshop on Recurring malware*, WORM’07, 2007.
- [16] M. G. Kang, H. Yin, S. Hanna, S. McCamant, and D. Song. Emulating emulation-resistant malware. In *VMSec’09: Proc. 1st ACM Workshop on Virtual Machine Security*, pages 11–22, 2009. URL <http://doi.acm.org/10.1145/1655148.1655151>.
- [17] N. Kisserli, J. Cappaert, and B. Preneel. Software security through targeted diversification. In *Proceedings Third Int. Workshop on Code Based Software Security Assessments*, 2007.
- [18] T. Kojm. Clam antivirus (ClamAV). page <http://clamav.net>, 2012.
- [19] C. Kolbitsch, E. Kirda, and C. Kruegel. The power of procrastination: detection and mitigation of execution-stalling malicious code. In *ACM Conference on Computer and Communications Security*, pages 285–296, 2011.
- [20] C. Lattner and V. Adve. LLVM: A compilation framework for lifelong program analysis & transformation. In *CGO’04: Proc. 2004 Int. Symp. Code Generation and Optimization*, 2004.
- [21] B. Lau and V. Svajcer. Measuring virtual machine detection in malware using dsd tracer. *Journal in Computer Virology*, pages 181–195, 2010.
- [22] M. Lindorfer, C. Kolbitsch, and P. Milani Comparetti. Detecting Environment-Sensitive Malware. In *Recent Advances in Intrusion Detection (RAID) Symposium*, 2011.
- [23] L. Martignoni, M. Christodorescu, and S. Jha. Omniunpack: Fast, generic, and safe unpacking of malware. In *In Proceedings of the Annual Computer Security Applications Conference (ACSAC)*, 2007.
- [24] J. Oberheide, M. Bailey, and F. Jahanian. PolyPack: an automated online packing service for optimal antivirus evasion. In *Proceedings of the 3rd USENIX conference on Offensive technologies*, WOOT’09, 2009.
- [25] R. Paleari, L. Martignoni, G. Fresi, and R. D. Bruschi. A fistful of red-pills: How to automatically generate procedures to detect CPU emulators. In *WOOT’09: Proc. USENIX Workshop on Offensive Technologies*, 2009.
- [26] R. Perdisci, A. Lanzi, and W. Lee. Classification of packed executables for accurate computer virus detection. *Pattern Recogn. Lett.*, 29(14), Oct. 2008.
- [27] T. Raffetseder, C. Krügel, and E. Kirda. Detecting system emulators. In *ISC’07: Int. Conf. Information Security*, pages 1–18, 2007.
- [28] P. Royal, M. Halpin, D. Dagon, R. Edmonds, and W. Lee. PolyUnpack: Automating the hidden-code extraction of unpack-executing malware. In *Proceedings of the 22nd Annual Computer Security Applications Conference*, ACSAC’06, 2006.
- [29] J. Rutkowska. Red pill... or how to detect VMM using (almost) one CPU instruction. <http://invisiblethings.org/papers/redpill.html>, 2004.